



DATABRICKS



LAKEBASE

Lakebase 101

Operational Postgres on the Lakehouse – a shared, governed memory layer for Databricks & UiPath agents.

Suncorp AI CoE

CDP

Infosec

60 minutes

■ AGENDA · 60 MINUTES

What we will cover

- ◆ **Context** – why this session, and what problem Lakebase solves for claims
- ◆ **Lakebase 101** – what it is, the autoscaling model, and where it fits
- ◆ **Reference architecture** – data flow: Source → Delta → Lakebase
- ◆ **Security & connectivity** – the identity chain and the UC / Postgres RBAC boundary
- ◆ **Lakeflow cost controls** – sync schedule, timeseries key, cost tagging
- ◆ **Live demo** – Data API vs direct SQL, and role vs ownership

SECTION 01

01

Context – why this session

■ CONTEXT

Why Lakebase, why now

- ◆ Agentic workflows need **low-latency operational state**, not just analytics
- ◆ In claims, **Databricks agents and UiPath agents must share the same claim state**
- ◆ Lakebase gives **Postgres speed** with **Unity Catalog governance**
- ◆ Goal: align **AI CoE, CDP and Infosec** on one safe reference pattern
- ◆ Outcome: a reviewed blueprint you can extend toward production

SECTION 02

02

Lakebase 101



■ LAKEBASE 101

What is Lakebase

- ◆ Fully-managed **PostgreSQL (OLTP)** built into Databricks
- ◆ Serverless **Autoscaling: 0.5–32 CU**, scales to zero when idle
- ◆ **Git-style branching** for instant dev / test copies
- ◆ Unity Catalog integration over the **open Postgres protocol**
- ◆ **Millisecond** reads and writes for apps, agents and serving

The autoscaling model & where it fits

Resource model

- Project › Branch › Endpoint
- Read-write + read-only endpoints
- Scale-to-zero — pay for what you use
- Branch from a point in time

Design principle

- An **acceleration / serving** layer
- **Not** the system of record
- Delta stays the source of truth
- A low-latency mirror of governed data

How autoscaling works

How it scales

- Compute in **Compute Units – 2 GB RAM each**, up to 32 CU
- Watches **CPU, memory** and **working-set size** (hot data)
- Adjusts within your range with **no database restart**
- Proactive – heads off slowdowns and out-of-memory

Guardrails you set

- **Min CU** – responsiveness floor at low traffic
- **Max CU** – cost ceiling for spikes
- Max – min spread is bounded (keeps scaling predictable)
- **Scale-to-zero** on an idle timeout – ~70% savings on bursty/dev

You set the floor and ceiling; Lakebase moves between them automatically on real load – no manual resizing, no restarts within the range.

Branching & budget controls

Branch like git

- Copy-on-write clones of prod in seconds
- Isolated dev, test & preview
- Carries schema, data, RPCs — a full clone
- Drop when done — no cleanup drift

Costs stay controlled

- No compute until an endpoint is attached
- Cap CU + scale-to-zero per endpoint
- Gate who can create via project ACLs
- Tag + budget policy for chargeback

Let teams self-serve branches without risking production or the bill — we'll see this open the live demo.


SECTION 03

03

Reference architecture – data flow

■ REFERENCE ARCHITECTURE

Source → Delta → Lakebase



Synced tables use a **primary key + timeseries key** to keep the latest row per key. Optional write-back flows agent state from Lakebase back to Delta for analytics.

■ REFERENCE ARCHITECTURE

Acceleration layer, not system of record

Lakebase IS

- A low-latency serving store
- Operational agent memory
- App state, features, sessions
- A governed mirror of Delta

Lakebase is NOT

- The source of truth
- A warehouse for BI / history
- Where you retain long-term data
- A bypass around UC governance


SECTION 04

04

Security & connectivity

■ SECURITY & CONNECTIVITY

Identity chain (documented — not shown live)



Follows the Suncorp identity design doc. Presented as reference architecture; the live demo uses the field-eng workspace.

■ SECURITY & CONNECTIVITY

UC / Postgres RBAC boundary

Unity Catalog plane

- Catalogs, schemas, lineage
- CDF history tables
- Databricks grants + audit
- Governs Delta & CDF outputs

Postgres RBAC plane

- Roles, RLS, GRANT / REVOKE
- Security-definer RPCs
- Data API role mapping
- Governs live Lakebase rows

Two planes, one identity – the same principal is governed differently on each side of the boundary.

Role vs ownership

Authenticated principal

- The SP / user identity
- Mapped to a Postgres role
- Least privilege: EXECUTE only
- Recorded in every audit row

Object owner

- Owns the tables and RPCs
- SECURITY DEFINER bypasses RLS
- Runtime never owns tables
- Data API needs
databricks_create_role

Identity for the Data API & attribution

Service principal per runtime — Data API

- The Data API maps the token to the **individual** identity
- Give each runtime its own SP role(`databricks_create_role`)
- Audit principal = the SP — precise per-runtime attribution
- The path the agents use in this demo

Group role — direct connection

- Map a group with `databricks_create_role(name, 'GROUP')`
- Members connect **as the group role** with their own token
- Not used by the Data API (it keys off the individual)
- Audit principal = the group — pass an actor id for per-person

Audit records the **authenticated principal** and the caller-supplied **logical agent id** on every operation — SPs give per-runtime attribution; a shared group role attributes to the group unless the caller identity is passed in.

SECTION 05

05

Lakeflow cost controls



LAKEFLOW COST CONTROLS

Sync policy & cost attribution

Sync policy + cadence

- SNAPSHOT / TRIGGERED / CONTINUOUS modes
- Continuous = freshest, highest cost
- Triggered or scheduled for batch freshness
- Right-size cadence to the freshness need

Cost attribution

- Tag jobs, warehouses & Lakebase
- Chargeback by team / use case
- **Scale-to-zero** cuts idle spend
- Monitor via UC system tables

■ OPERATING MODEL

Governed self-service branching

Platform team owns the policy

- Define one branch template — capped CU, scale-to-zero
- Run branch creation via a service principal
- Attach a budget policy and cost tags
- Monitor CU-hours in Unity Catalog system tables

Data teams self-serve

- Request a branch through automation, not tickets
- Receive an isolated clone of production
- Compute is bounded and scales to zero
- Branch expires automatically when work ends

Recommendation. Provision branches through a platform-owned automation that applies a governed template — data teams self-serve while the platform team controls compute size, autoscaling, and lifecycle centrally.

■ CHANGE MANAGEMENT

Shipping schema changes: the declarative line

Above the line — `bundle deploy`

- Project, branches, endpoints, roles, database
- Synced tables (Databricks derives the schema)
- Declared as `postgres_*` resources
- Reconciled to match the YAML

Below the line — `bundle run`

- App tables, columns, indexes, grants
- Ordered, idempotent `V1__` / `V2__` migrations
- Applied by a job, tracked in `schema_version`
- Promoted by switching target: `dev` → `stage` → `prod`

Recommendation. Ship Lakebase changes with Declarative Automation

Bundles — infrastructure as `postgres_*` resources, schema as versioned migrations, promoted by target; branch to validate a change, then promote the migration file, not the branch.

The declarative line — a simple example

You author the columns → below

- An orders table your app owns
- You write the `CREATE TABLE` in a migration (`V1_...sql`)
- Applied by `bundle run`

Databricks derives the columns → above

- A `customer_segments` table synced from a Delta table
- You name the **source + key**, never the columns
- Declared as `postgres_synced_tables; bundle deploy`

Rule of thumb. If you write the columns, it ships as a migration (below the line). If Databricks derives or manages them, it is a `postgres_*` resource (above the line).

SECTION 06 · LIVE

06

Live demo — field-eng workspace

■ [LIVE DEMO · START HERE](#)

Branch production — clone it, safely

- ◆ Create a dev branch from production — **instant, copy-on-write**
- ◆ It carries the same schema, RPCs and agent memory — a full point-in-time clone
- ◆ Attach a **budget-safe** endpoint — small CU cap + scale-to-zero
- ◆ Experiment on dev; **production is untouched** — true isolation
- ◆ Throw the branch away when done — storage only until an endpoint runs

Verified live: dev cloned 5 tables + 4 RPCs + the seeded memory; a write on dev left production at 2 rows.

■ LIVE DEMO

Suncorp claims-center · agent memory

- ◆ Two agents share claim-scoped memory: `databricks.claim_triage.v1` & `uipath.document_review.v1`
- ◆ Runtime identities call **approved RPCs** – never touching tables directly
- ◆ Every operation writes an **append-only audit event** (reads included)
- ◆ **CDF** exposes row changes as Unity Catalog Delta history
- ◆ **Cross-claim access is denied** – boundaries fail closed

■ LIVE DEMO

Data API vs direct SQL

Data API (PostgREST)

- HTTPS + OAuth bearer token
- RPC + CRUD, **no Postgres driver**
- Ideal for UiPath / low-code
- Role via databricks_create_role

Direct SQL

- psycopg / psycopg2 + OAuth cred
- Full SQL and transactions
- Apps, notebooks, migrations
- Same roles, RLS and RPCs

■ LIVE DEMO · VERIFIED

What the demo proves

CHECK	PATH	RESULT
Read / write / task RPCs	Data API · as service principal	● 200 OK
Cross-claim completion	Data API · as service principal	● 403 · enforced
Direct table access	Data API · as service principal	● 403 · enforced
Principal vs logical agent	Append-only audit event	● both recorded
Row-change history	Unity Catalog CDF	● queryable

Green means the boundary behaved correctly – including the two 403s, where denial is the intended, verified outcome.

■ WRAP-UP

Security takeaways

- ◆ Runtime principals get `EXECUTE` on approved RPCs only – **no table DML**
- ◆ RLS is enabled **fail-closed**; direct access returns nothing or is denied
- ◆ Security-definer RPCs enforce explicit **claim + business-unit** checks
- ◆ Audit records both the **authenticated principal** and the **logical agent id**
- ◆ CDF gives immutable **write history** in UC; reads are audited separately

■ WRAP-UP

Recommendations & next steps

- ◆ Use **separate service principals per agent runtime** for strong attribution
- ◆ Keep Lakebase as the **serving layer**; Delta remains the system of record
- ◆ Standardize the identity chain with **Infosec** (Entra → OAuth → PG role)
- ◆ Set sync cadence & cost tags with **CDP** for freshness and chargeback
- ◆ Review the SQL + RPC pattern with **security** before production

THANK YOU

Questions & discussion

Lakebase 101 · Suncorp AI CoE, CDP & Infosec – the claims-center agent-memory demo is a synthetic, non-production starter, reviewed end-to-end on the field-eng workspace.
